

# Concurrency Control in Transactional Systems: Autumn 2019

## Programming Assignment 2: Implementing BTO and MVTO algorithms

Submission Date: 10th November 2019, 9:00 pm

**Goal:** The goal of this assignment is to implement BTO and MVTO algorithms studied in the class. Implement both these algorithms in C++.

**Details.** As shown in the book, you have to implement BTO and MVTO algorithms studied in the class in C++. But unlike the version that we studied in the class, you have to implement these algorithms using optimistic concurrency control approach, i.e. all writes become visible only after commit. It can be seen that the advantage of this approach: upon an abort of a transaction, no rollback is necessary as none of the writes of the transactions will ever be visible.

**MVTO Variants:** An important requirement with MVTO algorithm is to delete the unwanted (garbage) versions. There are two ways to do it. Either through garbage collection. Or have a fixed number of versions, say like 5. Once all the 5 versions have been written onto, the next transaction write overwrites the first version. This way at any time only a fixed number of versions (5 in this case) will be in the system for any data-item.

So, you have to implement the garbage collection procedure as well that we studied in the class. So for MVTO, you have to implement three variants:

1. MVTO: Normal MVTO that we studied in the class.
2. MVTO-gc: The variant of MVTO with garbage collection.
3. K-MVTO: The variants with k versions of a data-item maintained any time.

**Methods to Implement:** Just like in previous assignment, you have to implement the following methods for both the algorithms:

- *begin\_trans()*: It begins a transactions and returns a unique transaction id, say  $i$
- *read( $i, x, l$ )*: Transaction  $t_i$  reads data-item  $x$  into the local value  $l$ .
- *write( $i, x, l$ )*: Transaction  $t_i$  writes to data-item  $x$  with local value  $l$ .
- *tryC( $i$ )*: Transaction  $t_i$  wants to commit. The return of this function is either  $a$  for abort or  $c$  for commit.

To test the performance of both the algorithms, develop an application, opt-test which you developed in the previous assignment. It is as follows: Once, the program starts, it creates  $n$  threads and an array of  $m$  shared variables. Each of these threads, will update the shared array randomly. Since the threads could simultaneously update the shared variables of the array, the access to shared variables have to be synchronized. The synchronization is performed using the above mentioned methods of BTO & MVTO.

**Test Program:** To test your libraries, you have to develop a test program, say opt-test like in the previous assignment. The pseudocode opt-test given is as follows:

Listing 1: main thread

```

1 void main()
2 {
3     ...
4     ...
5     // create a shared array of size m
6     shared [] = new SharedArray[m];
7     ...
8     ...
9     create n updtMem threads;
10 }
```

Listing 2: updtMem thread

```

1
2 void updtMem()
3 {
4     int status = abort;           // declare status variable
5     int abortCnt = -1;           // keeps track of abort count
6
7     long critStartTime , critEndTime;
8
9     critStartTime = getSysTime(); // keep track of critical section start time
10
11     // getRand(k) function used in this loop generates a random number in the range 0 .. k
12     do
13     {
14         id = begin_trans(); // begins a new transaction id
15         randIters = getRand(m); // gets the number of iterations to be updated
16
17         int locVal;
18         for (int i=0; i<randIters; i++)
19         {
20             // gets the next random index to be updated
21             randInd = getRand(m);
22
23             // gets a random value using the constant constVal
24             randVal = getRand(constVal);
25
26             // reads the shared value at index randInd into locVal
27             read(id , shared[randInd] , locVal);
28
29             logfile << "Thread id " << pthread_self() << "Transaction " << id <<
30             " reads from" << randInd << " a value " << locVal << " at time " <<
31             getSysTime;
32 }
```

```

33         // update the value
34         locVal += randVal;
35
36         // request to write back to the shared memory
37         write(id, shared[randInd], locVal);
38
39         logFile << "Thread id " << pthread_self() << "Transaction " << id <<
40         " writes to " << randInd << " a value " << locVal << " at time " <<
41         getSysTime;
42
43         // sleep for a random amount of time which simulates some complex computation
44         randTime = getExpRand( $\lambda$ );
45         sleep(randTime);
46     }
47
48     status = tryCommit(id); // try to commit the transaction
49     logFile << "Transaction " << id << " tryCommits with result "
50     << status << " at time " << getSysTime;
51     abortCnt++; // Increment the abort count
52 }
53 while (status != commit);
54
55 critEndTime = getSysTime(); // keep track of critical section end time
56 }

```

Here *randTime* is an exponentially distributed with an average of  $\lambda$  mill-seconds. The objective of having this time delays is to simulate that these threads are performing some complicated time consuming tasks. It can be seen that the time taken by a transaction to commit, *commitDelay* is defined as *critEndTime* – *critStartTime*.

**Input:** The input to the program will be a file, named inp-params.txt, consisting of all the parameters described above:  $n, m, constVal, \lambda$ . A sample input file is: 10 10 100 20.

**Output:** Your program should output to a file in the format given in the pseudocode for each algorithm. A sample output is as follows:

```

Thread 1 Transaction 1 reads 5 a value 0 at time 10:00
Thread 2 Transaction 2 reads 7 a value 0 at time 10:02
Thread 1 Transaction 1 writes 5 a value 15 at time 10:05
Thread 2 Transaction 2 tryCommits with result abort at time 10:10

```

.  
.  
.

The output is essentially a history. By inspecting the output one should be able to verify the serializability of your implementations.

**Report:** You have to submit a report for this assignment. This report should contain a comparison of the performance of BTO & MVTO's variants. The comparison must consist of two graphs like in the previous assignment:

- A graph comparing the average time taken by a transaction to successfully commit, i.e. *commitDelay* in BTO & MVTO. The x-axis should the number of threads while the y-axis should be the average of *commitDelays*. This graph will have four curves: (1) BTO (2) MVTO (3) MVTO-gc (4) K-MVTO.

- A graph comparing the average abort count, i.e. the number of times a transaction abort before it can successfully commit . The x-axis should the number of threads while the y-axis should be abort count. Thus, this graph similar to the previous graph will have four curves: (1) BTO (2) MVTO (3) MVTO-gc (4) K-MVTO.

You must run these algorithms multiple times to obtain performances values. You run these algorithms varying the number of threads from 10 to 50 while keeping other parameters same. Please have  $m$  to be fixed at 10 in all these experiments. Finally in your report, you must also give an analysis of the results while explaining any anomalies observed.

**Deliverables:** You have to submit the following:

- The source file containing the actual program to execute
- A readme.txt that explains how to execute the program
- The report as explained above

Zip all the three files and name it as ProgAssn2-BTO\_MVTO-<rollno>.zip. Then upload it on the google classroom page of this course by the above mentioned deadline.